

SentiSum Product Specialist Interview Preparation Guide

Complete preparation materials for Bhavya Mahajan

Table of Contents

1. Introduction & Interview Overview
2. AI Fundamentals (Explained A-Z)
3. Assignment Deep Dive Q&A
4. Product Management Frameworks
5. KPI Tree for FinServeCo
6. Practice Questions & Answers
7. Final Preparation Checklist

1. Introduction & Interview Overview

About the Role

Position: Product Specialist - AI

Company: SentiSum

Interview Format: 90-minute discussion

Focus: Assignment walkthrough + deep dive questions

What SentiSum Does

SentiSum is an **AI-native Voice of Customer (VoC) platform** that unifies customer feedback and drives retention outcomes. Key features:

- **Kyo AI Assistant:** Acts as "Customer Second Brain"
- **Omnichannel Analysis:** Support tickets, reviews, calls, surveys
- **Real-time Insights:** Anomaly detection and alerting
- **Root Cause Analytics:** Separates product vs process vs agent issues

Interview Structure

Part 1 (30-40 min): Walk through your assignment

Part 2 (25-30 min): Deep dive technical and product questions

Part 3 (15-20 min): Product sense scenarios

Part 4 (5-10 min): Your questions for them

2. AI Fundamentals (Explained A-Z)

2.1 Core Concepts

What is a Large Language Model (LLM)?

Simple Definition: An AI model trained on massive amounts of text that can understand and generate human-like language.

Technical Detail: Neural networks with billions of parameters trained using transformer architecture on diverse text datasets. They learn patterns, context, and relationships in language.

Examples: GPT-4, Claude, PaLM, LLaMA

Application: Your CoT prompts are processed by an LLM to analyze support conversations and generate structured JSON outputs.

What is Prompt Engineering?

Simple Definition: The art and science of crafting instructions (prompts) to get desired outputs from AI models.

Why It Matters:

- Guides AI to understand your intent
- Determines quality and relevance of outputs
- Critical skill for working with LLMs

Best Practices:

1. Be specific and clear
2. Provide context
3. Use examples (few-shot)
4. Iterate and refine
5. Test with edge cases

Your Application: Your entire Part 1 is prompt engineering - you designed how AI should evaluate conversations systematically.

Chain-of-Thought (CoT) Prompting

Simple Definition: Asking AI to show its step-by-step reasoning before giving a final answer, like showing your work in math class.

Why It Works:

- Breaks down complex problems
- Improves accuracy on reasoning tasks
- Increases transparency
- Reduces hallucinations

Types of CoT:

1. **Zero-Shot CoT:** Simply add "Let's think step by step" to your prompt
 - Quick to implement
 - No examples needed
 - Less reliable for complex tasks
2. **Manual CoT:** Provide hand-crafted examples showing reasoning
 - High accuracy
 - Time-consuming to create
 - Hard to scale
3. **Auto-CoT:** AI generates its own reasoning examples
 - Scalable approach
 - Learns from patterns
 - Requires validation

Your Implementation:

You used CoT to make AI analyze each quality dimension step-by-step like a human quality auditor would, ensuring structured and explainable evaluations.

2.2 LLM Evaluation & Validation

What is LLM Evaluation?

Definition: The process of testing and measuring how well LLMs perform in real-world situations.

Why Critical:

- Ensures accuracy before production
- Identifies weaknesses
- Guides improvement
- Builds stakeholder confidence

Key Metrics:

1. **Answer Relevancy:** Does the output address the input?
2. **Correctness:** Is the output factually accurate?
3. **Hallucination:** Does it contain made-up information?

4. **Contextual Relevancy:** Is the retrieved context appropriate?
5. **Task Completion:** Did it achieve the goal?

LLM-as-a-Judge

Definition: Using one LLM to evaluate outputs from another LLM (or itself).

Advantages over Traditional Metrics (BLEU, ROUGE):

- Captures semantic nuance
- Understands context
- Evaluates quality holistically
- More aligned with human judgment

Techniques:

- G-Eval: Uses natural language rubrics
- Pairwise comparison
- Reference-based scoring

Your Application:

Your validation framework uses LLM-as-a-Judge in the second layer, where AI evaluations are compared against human ground truth.

Model Validation Techniques

1. Train/Test Split

- Divide data into training (70-80%) and testing (20-30%)
- Train on one set, evaluate on unseen set
- Simple but can be unstable

2. Cross-Validation

- K-Fold: Split data into K parts, rotate which is test set
- More robust than single split
- Reduces variance in performance estimates

3. Human-in-the-Loop

- Expert evaluators establish ground truth
- Critical for subjective tasks
- Expensive but necessary

Your Two-Layer Validation:

Layer 1: Human-in-the-Chain

- 3-5 expert evaluators
- Label 100-200 conversations
- Measure inter-rater reliability
- Create gold standard dataset

Layer 2: LLM-as-Evaluator

- Compare AI predictions to human labels
- Calculate agreement rate (target: ≥85%)
- Measure precision and recall per dimension
- Identify systematic errors for prompt refinement

Understanding Hallucination

Definition: When AI generates false or made-up information confidently.

Why It Happens:

- Model predicts plausible-sounding text
- No grounding in real data
- Fills gaps with "reasonable" guesses

How to Prevent:

1. Provide clear context and constraints
2. Use few-shot examples with correct patterns
3. Request evidence/quotes from source
4. Implement confidence thresholds
5. Add human review for critical outputs

Your Risk Mitigation:

- CoT prompting reduces hallucination
- Evidence field in JSON (quotes from transcript)
- Confidence scores for each prediction
- Human validation layer catches errors

Precision vs Recall Trade-off

Precision: Of all positive predictions, how many are correct?

- Formula: $\text{True Positives} / (\text{True Positives} + \text{False Positives})$
- High precision = Few false alarms

Recall: Of all actual positives, how many did we catch?

- Formula: $\text{True Positives} / (\text{True Positives} + \text{False Negatives})$
- High recall = Catch all important cases

The Trade-off:

- Increasing one often decreases the other
- Optimize based on business needs

Your Application:

For **policy violations** (critical alert):

- Prioritize RECALL - can't miss compliance issues
- Accept some false positives for human review

For **empathy coaching** (non-critical):

- Balance precision and recall
- Use F1 score for optimization

2.3 Key AI Terminology

Embedding: Converting text into numerical vectors that capture semantic meaning. Used in your friction themes clustering.

Sentiment Analysis: Determining emotional tone (positive/negative/neutral). Your "Emotion Management" dimension tracks sentiment change.

Anomaly Detection: Identifying unusual patterns. Your "Post-Release Spike" detection catches issues early.

Model Drift: Performance degrading over time as real data changes. Your guardrail metrics monitor this.

Fine-tuning: Training a pre-trained model further on domain-specific data. Potential future enhancement for FinServeCo.

Confidence Score: AI's self-reported certainty (0.0-1.0). Included in your JSON for each dimension.

Ground Truth: Verified correct answers for measuring performance. Your gold standard examples and human labels.

F1 Score: Harmonic mean of precision and recall. Single metric for balanced evaluation.

Context Window: Maximum text length AI can process at once. Important for long conversations.

Contextual Tagging: Understanding context to apply correct labels. Your topic/intent/root cause classification.

3. Assignment Deep Dive Q&A

Part 1: Problem Structuring & Prompt Frameworks

Q: "Walk me through your approach to defining conversation quality dimensions"

Your Answer:

1. Started with Problem Framework

- Identified key actors: Customers, CX Managers, Product/Ops Teams, Business Leadership
- Evidence: High repeat contacts, inconsistent tagging, leadership frustration
- Why now: B2B churn risk, rising BPO costs, lack of rapid root cause detection

2. Defined 6 Quality Dimensions

- Each maps to specific FinServeCo pain point
- Multi-dimensional vs single score approach
- Stakeholder-specific alerts

3. Chose CoT Prompting

- Makes AI think like quality auditor
- Step-by-step, transparent, explainable
- Reduces hallucination risk

4. Designed JSON Schema

- Standardized across channels
- Machine-readable for automation
- Supports real-time alerting

Follow-up Points:

- Why 6 dimensions specifically? Each addresses a different root cause
- Considered 8-10 initially, consolidated to essential ones
- Can expand later based on pilot learnings

Q: "Why Chain-of-Thought prompting specifically?"

Your Answer:

Analogy: Like showing your work in math class - you don't just write the answer, you show each step.

Why CoT > Alternatives:

vs Zero-Shot:

- Zero-shot: "Analyze this conversation"

- CoT: "Step 1: Check for regulatory disclosures. Step 2: Verify calculations..."
- Result: More accurate, explainable

vs **Few-Shot**:

- Few-shot helps but doesn't show reasoning
- CoT shows WHY model made decision
- Critical for building trust with stakeholders

Benefits for This Use Case:

1. **Transparency**: CX managers see exact reasoning
2. **Debuggability**: Can identify where model went wrong
3. **Consistency**: Structured approach reduces variance
4. **Auditability**: Important for compliance (fintech context)

Q: "How does your JSON schema enable scalability?"

Your Answer:

Standardization:

- Same structure for chat, email, voice (after transcription)
- Easy to add new channels without redesign
- Consistent querying across all data

Automation:

- Machine-readable format
- Direct integration with CRM/support tools
- API-ready for downstream systems

Real-time Processing:

- Streaming JSON as conversations complete
- Immediate alert triggering
- No manual processing bottleneck

Analytics:

- Structured fields enable SQL queries
- Time-series analysis on any dimension
- Cross-channel correlation studies

Extensibility:

- Can add new dimensions without breaking schema

- Backward compatible
- Version control friendly

Q: "What's the difference between your approach vs traditional CSAT surveys?"

Your Answer:

Aspect	Traditional CSAT	Your AI Framework
When	After interaction (lagging)	Real-time analysis (leading)
Coverage	5-30% response rate	100% of conversations
Depth	Single number (1-5)	6 dimensions + root cause
Actionability	"Customer unhappy"	"Agent missed policy X; coach on Y"
Stakeholders	CX only	CX, Product, Ops, Legal - each gets relevant alerts
Prevention	Reactive	Proactive (catches issues early)
Cost	Survey fatigue, incentives	Automated, no customer burden

Key Insight: CSAT tells you WHAT (satisfaction level). Your framework tells you WHY (root causes) and HOW TO FIX (targeted actions).

Part 2: Insights & Analytics Thinking

Q: "Explain your Post-Release Spike detection methodology"

Your Answer:

The Problem:

- New feature releases can introduce bugs
- Support tickets spike within hours
- Current systems detect too late

Your Solution:

Step 1: Establish Baseline

- Track ticket volume per topic (7-day moving average)
- Calculate normal variance range
- Account for day-of-week patterns

Step 2: Real-time Monitoring

- Correlate with product release calendar
- Flag spikes > 2 standard deviations

- Classify severity: magnitude × acceleration

Step 3: Root Cause Attribution

- Link spike to specific feature/release
- Identify if product bug vs process gap
- Extract common error patterns

Step 4: Alert & Action

- Critical: Immediate Slack alert to Product + Ops
- Include: Spike timeline, affected customers, error samples
- Enable: Quick rollback decision or workaround

Success Metrics:

- Detection within 2 hours of spike start
- False positive rate < 5%
- Mean time to resolution reduced 30%

Q: "How do your insights drive cross-functional alignment?"

Your Answer:

The Challenge: Different teams need different insights at different times.

Your Solution: Stakeholder-Specific Design

For CX Managers:

- Daily empathy trend heatmap
- Agent-level coaching recommendations
- Threshold alerts for quality drops

For Product Team:

- Weekly friction themes report (bubble chart)
- Post-release spike alerts (real-time)
- Root cause impact analysis (monthly)

For Operations:

- Real-time compliance alerts
- Process gap identification
- FCR impact by issue type

For Business Leadership:

- Monthly executive summary

- Churn risk indicators
- ROI metrics (cost saved, satisfaction improved)

Alignment Mechanism:

- Shared data layer (everyone sees same source)
- Role-based views (relevant insights only)
- Cross-functional OKRs (insights → actions → outcomes)

Part 3: Operational Planning & Communication

Q: "Walk me through your 3-month pilot structure"

Your Answer:

Phase 1: Foundation & Calibration (Weeks 1-3)

Goal: Establish baseline and build trust

Milestones:

- Stakeholder kickoff and alignment
- Secure data access (historical tickets, CSAT, CRM)
- Define 20 gold standard examples per key issue
- Build initial prompt library
- Establish baseline metrics

Deliverables:

- Project charter with success criteria
- Data pipeline and credentials
- Prompt library v1.0
- Baseline performance report

Why 3 weeks: Need time for data access approvals, SME availability for gold standard labeling, and baseline stabilization.

Phase 2: Model Development & Validation (Weeks 4-6)

Goal: Build and validate AI quality model

Milestones:

- Train conversation quality model
- Build topic/intent classification
- Develop root cause detection

- Calibrate alert thresholds
- Achieve $\geq 85\%$ agreement with humans

Deliverables:

- Validated AI model
- Topic taxonomy & classification system
- Root cause detection module
- Threshold configuration document

Risk Mitigation:

- Weekly model performance reviews
- Early escalation if accuracy < 80%
- SME involvement throughout

Phase 3: Dashboard & Integration (Weeks 7-9)

Goal: Make insights accessible and actionable

Milestones:

- Build CX Manager coaching dashboard
- Build Product insights dashboard
- Build Operations compliance alerts
- Integrate with Zendesk/Slack
- User acceptance testing

Deliverables:

- 3 role-specific dashboards
- Slack alert integration
- API endpoints for data export
- User training documentation

Success Criteria:

- Dashboard loads < 3 seconds
- Zero critical bugs in UAT
- Training completion by pilot teams

Phase 4: Pilot Launch & Iteration (Weeks 10-12)

Goal: Learn fast, improve fast

Milestones:

- Launch with 2-3 support teams (100-200 agents)
- Daily monitoring and issue triage
- Collect user feedback (weekly surveys)
- Measure adoption and impact KPIs
- Refine based on live data

Deliverables:

- Pilot performance report
- User feedback synthesis
- Updated model/prompt configuration
- Rollout readiness assessment

What Good Looks Like:

- 70%+ of managers using dashboards weekly
- 85%+ agreement with human evaluators
- 15% reduction in avg resolution time (even small progress)
- No critical incidents or data breaches

Q: "How do you balance daily execution with strategic communication?"

Your Answer:

Communication Rhythm = Audience + Cadence + Medium

Daily (Team-Facing):

- **Audience:** Core project team (you, eng, data science)
- **Medium:** Slack standup (async)
- **Content:** Blockers, quick wins, help needed
- **Time:** < 10 min to read/respond

Weekly (Cross-Functional):

- **Audience:** CX, Product, Ops leads
- **Medium:** Email + Dashboard snapshot
- **Content:**
 - Progress vs milestones
 - Metrics update
 - Top 3 insights from data
 - Blockers and decisions needed
- **Time:** 5 min read + 15 min discussion if needed

Monthly (Leadership):

- **Audience:** Execs, broader organization
- **Medium:** Recorded demo + written summary
- **Content:**
 - Business impact (cost, satisfaction, retention)
 - Wins and learnings
 - What's next
 - Invite feedback
- **Time:** 10 min video + 2-page doc

Key Principle: Push insights, pull decisions. Don't wait for people to ask - proactively share what's relevant to them.

Part 4: Validation & Metrics

Q: "Explain your two-layer validation approach"

Your Answer:

Why Two Layers?

Layer 1 alone (humans only): Doesn't scale, expensive

Layer 2 alone (AI only): Risk of systematic bias, no ground truth

Both together: Best of both worlds

Layer 1: Human-in-the-Chain (Establish Ground Truth)

Process:

1. Recruit 3-5 expert evaluators (CX managers, QA specialists)
2. Provide evaluation rubric and training
3. Each evaluator labels 100-200 conversations independently
4. Measure inter-rater reliability (Cohen's Kappa)
5. Resolve disagreements through consensus discussion
6. Create gold standard dataset

Metrics:

- Inter-rater agreement ≥ 0.75 (substantial agreement)
- Coverage: All major issue types represented
- Diversity: Different channels, agents, severity levels

Why Critical: Sets objective standard that reflects FinServeCo's unique quality definition, not generic AI assumptions.

Layer 2: LLM-as-Evaluator (Validate at Scale)

Process:

1. AI model analyzes same 100-200 conversations
2. Compare AI labels vs human ground truth
3. Calculate agreement rate per dimension
4. Identify systematic errors (e.g., AI always misses sarcasm)
5. Refine prompts and retry
6. Repeat until $\geq 85\%$ agreement achieved

Metrics:

- **Overall Agreement:** 85%+ target
- **Per-Dimension Precision/Recall:** ≥ 0.85 each
- **False Positive Rate:** $< 5\%$ for critical alerts
- **Consistency:** $\geq 90\%$ same answer on repeated runs

Continuous Improvement:

- Monthly re-validation with new conversations
- Track drift in agreement over time
- Retrain or update prompts as needed

Q: "How would you handle disagreement between AI and human evaluators?"

Your Answer:

Step 1: Categorize the Disagreement

Type A: Edge Cases

- Ambiguous conversations (e.g., polite but unhelpful)
- Lack of context in transcript
- Solution: Clarify rubric, add examples for edge cases

Type B: Systematic Errors

- AI consistently misses sarcasm
- AI over-penalizes long but informative responses
- Solution: Refine prompt logic, add specific instructions

Type C: Human Error

- Evaluator fatigue, misunderstanding rubric
- Inter-rater disagreement

- Solution: Re-train evaluators, improve rubric clarity

Step 2: Investigate Root Cause

- Pull 20 examples of disagreements
- Analyze patterns:
 - Is it specific to certain agents/topics/channels?
 - Is confidence score low for these cases?
 - Do multiple humans also disagree on these?

Step 3: Intervention

If AI Error:

- Add explicit instruction to prompt
- Include few-shot examples of correct handling
- Adjust threshold if over/under-sensitive
- Consider fine-tuning if systematic

If Human Error:

- Improve evaluation guidelines
- Add calibration sessions
- Increase number of raters for these cases

If Genuinely Ambiguous:

- Create "requires human judgment" category
- Set confidence threshold trigger
- Route to human review queue

Step 4: Document and Monitor

- Log all disagreements and resolutions
- Track disagreement rate over time (should decrease)
- Share learnings with team
- Update validation process

Success Metric: Disagreement rate < 15% and decreasing by 2% per month.

Q: "What metrics beyond CSAT would indicate quality improvement?"

Your Answer:

Leading Indicators (Predict Future Outcomes):

1. **First Contact Resolution (FCR)**

- % of issues resolved in first interaction
- Target: +10% improvement
- Why: Directly tied to conversation quality

2. Agent Response Quality Score

- Your 6-dimension average
- Target: 15-point increase from baseline
- Why: More granular than CSAT

3. Context Utilization Rate

- % of conversations where agent accessed history
- Target: From 40% → 65%
- Why: Reduces repeat questions, shows proactive behavior

4. Root Cause Capture Rate

- % of conversations with diagnostic data
- Target: From 50% → 75%
- Why: Enables Product team to fix systemic issues

Efficiency Indicators:

5. Average Resolution Time (ART)

- Hours from ticket open to close
- Target: -15% reduction
- Why: Better quality often means faster resolution

6. Repeat Contact Rate

- % customers contacting again within 7 days
- Target: -20% reduction
- Why: Clarity of resolution directly impacts this

Behavioral Indicators:

7. Empathy Phrase Usage

- Avg empathy indicators per conversation
- Target: +30% increase
- Why: Trainable behavior that impacts sentiment

8. Policy Compliance Rate

- % conversations with no violations
- Target: 100% for critical, 95% for important
- Why: Risk mitigation and legal safety

Business Impact (Lagging but Important):

9. Customer Effort Score (CES)

- "How easy was it to resolve your issue?"
- Target: Decrease from 4.2 → 3.5 (on 7-point scale)
- Why: Quality directly reduces customer effort

10. Net Retention Rate

- % customers retained - % lost
- Target: +5 percentage points
- Why: Ultimate business outcome

Relationship: Leading indicators improve → Lagging indicators improve → Business metrics improve (with 30-90 day lag)

4. Product Management Frameworks

4.1 KPI Trees

Definition: Hierarchical representation of KPIs showing how metrics contribute to business goals.

Structure:

- **Root KPIs:** Top-level business objectives (e.g., Customer Retention)
- **Branch KPIs:** Specific drivers (e.g., Support Effectiveness)
- **Leaf KPIs:** Operational metrics (e.g., FCR, ART)

How to Build:

1. Start with Core Objective

- What's the ultimate business goal?
- Must be measurable and time-bound

2. Decompose into Drivers

- Ask "What influences this?"
- Create branches for each driver

3. Define Leaf Metrics

- Actionable, measurable, assignable to teams
- Should be improvable through specific actions

4. Establish Relationships

- Quantify how each metric impacts parent
- Test assumptions with data

5. Assign Ownership

- Each metric has a DRI (Directly Responsible Individual)
- Clear accountability

Benefits:

- Aligns team efforts to business outcomes
- Identifies gaps in measurement
- Facilitates data-driven decisions
- Optimizes resource allocation

4.2 Metrics Hierarchy

Product Metrics Framework:

Level 0: North Star Metric

- Single most important metric
- Represents value to users and business
- Example: "Weekly Active Engaged Users"

Level 1: Primary Drivers

- 3-5 key metrics that drive North Star
- Example: Acquisition, Activation, Retention

Level 2: Component Metrics

- Specific levers for each driver
- Example: Conversion rate, Time to value, Churn rate

Leading vs Lagging:

- **Leading:** Predict future outcomes (e.g., quality scores)
- **Lagging:** Measure past results (e.g., CSAT, churn)
- Balance both in your hierarchy

4.3 User Research Methods

Qualitative Research (Understand WHY)

User Interviews:

- One-on-one conversations
- Deep insights into motivations
- When: Discovering problems, validating solutions

Usability Testing:

- Watch users interact with product
- Identify friction points
- When: Evaluating prototypes, finding UX issues

Focus Groups:

- Group discussions
- Diverse perspectives
- When: Exploring broad themes, early ideation

Quantitative Research (Measure WHAT and HOW MANY)

Surveys:

- Structured questions
- Large sample sizes
- When: Validating hypotheses, prioritizing features

A/B Testing:

- Compare two versions
- Statistical significance
- When: Optimizing existing features

Analytics:

- Behavioral data
- Usage patterns
- When: Understanding actual usage, finding drop-offs

When to Use Each:

- **Generative (early):** Qualitative to discover problems
- **Evaluative (later):** Quantitative to test solutions
- **Best Practice:** Combine both for comprehensive understanding

4.4 Hypothesis-Driven Development

Structure:

"We believe we will achieve **[business outcome]**
if **[persona]**
solves their need in **[user outcome]**
using **[feature]**."

Example from Your Assignment:

"We believe we will achieve a **15% reduction in repeat contacts** if **CX managers** solve their need in **identifying coaching gaps** using **dimension-specific quality dashboards**."

Validation Process:

1. Define Hypothesis

- Clear, testable, specific

2. Design Experiment

- What will you build/test?
- How will you measure?

3. Set Success Metrics

- Quantitative threshold for success
- Timeline for evaluation

4. Execute & Learn

- Build minimum viable solution
- Collect data
- Analyze results

5. Iterate or Pivot

- Validated: Invest more
- Invalidated: Learn and adjust

Benefits:

- Reduces risk of building wrong thing
- Forces clear thinking about value
- Creates learning culture
- Connects features to outcomes

4.5 OKRs (Objectives & Key Results)

Framework:

Objective: Qualitative, inspirational goal

Key Results: Quantitative measures of success (2-5 per objective)

Example for Your Project:

Objective: Transform FinServeCo's support from reactive to proactive

Key Results:

1. Achieve 85% agreement rate with human quality evaluators
2. Reduce average resolution time by 15%
3. Increase First Contact Resolution rate by 10%
4. Deploy to 100+ agents with 70%+ dashboard adoption

OKRs vs Roadmaps:

- **OKRs:** OUTCOMES (what to achieve)
- **Roadmaps:** OUTPUTS (what to deliver and when)
- OKRs guide strategy, Roadmaps show execution

Best Practices:

- Set ambitious but achievable (70% success rate)
- Time-bound (usually quarterly)
- Transparent across organization
- Regularly review and adjust

4.6 UX/UI Design Principles

Core Pillars:

1. Usability

- Easy to learn
- Efficient to use
- Memorable
- Error-tolerant

2. Consistency

- Visual consistency (colors, fonts, spacing)
- Functional consistency (similar actions work similarly)
- Platform consistency (follow OS conventions)

3. Clarity

- Clear labels and instructions
- Obvious next steps
- No ambiguity

4. Feedback

- System responds to actions
- Progress indicators
- Error messages are helpful

5. Accessibility

- Works for users with disabilities
- Keyboard navigation
- Screen reader compatible

Application to Your Dashboards:

- Consistent color coding (red=critical, yellow=warning)
- Clear metric labels (not just "Score 72")
- Hover tooltips for definitions
- Responsive design for different screens

5. KPI Tree for FinServeCo

Complete Hierarchy (19 Metrics)

Level 0: NORTH STAR

└ Customer Retention Rate

Level 1: PRIMARY DRIVERS

└ Customer Satisfaction (CSAT)

└ Product Value Delivered

└ Support Effectiveness ★ [YOUR FOCUS]

Level 2: SUPPORT QUALITY METRICS

└ Conversation Quality Score ★★ [YOUR AI FRAMEWORK]

└ First Contact Resolution (FCR)

└ Average Resolution Time (ART)

└ Repeat Contact Rate

Level 3: QUALITY DIMENSIONS (Your 6)

└ Accuracy & Correctness (<70% = CRITICAL)

└ Context Awareness (<40% = URGENT)

└ Root Cause Depth (<50% = Product Alert)

└ Clarity of Resolution (<50% = Operations Alert)

└ Efficiency (<30% = Coaching Alert)

└ Emotion Management (0% = Churn Risk)

Level 4: ACTIONABLE SUB-METRICS

└ Product Issues Identified

└ Process Gaps Identified

└ Policy Violations

└ CRM Utilization Rate

└ Empathy Phrase Usage

How to Present This in Interview

Start at the Top:

"FinServeCo's ultimate goal is Customer Retention. That's the North Star."

Trace Downward:

"Retention is driven by three things: satisfaction, product value, and support effectiveness. My assignment focused on the support branch."

Show Your Contribution:

"Within support effectiveness, I created the Conversation Quality Score as a leading indicator. It's composed of 6 dimensions, each measuring a different aspect of quality."

Connect to Leaf Metrics:

"Each dimension breaks down into specific, actionable metrics. For example, Root Cause Depth tracks how many product issues we identify - directly helping the Product team prioritize fixes."

Stakeholder Mapping:

"Different levels serve different stakeholders:

- Level 0-1: Business leadership
- Level 2-3: CX and Product managers
- Level 4: Frontline agents and analysts"

6. Practice Questions & Answers

General Product Questions

Q: "How would you prioritize feature requests for this product?"

Framework: ICE Score

Impact: How much will this improve key metrics? (1-10)

Confidence: How certain are we this will work? (%)

Effort: How much work required? (1-10, inverted)

Score = (Impact × Confidence) / Effort

Your Answer:

"For the AI Quality Framework, I'd prioritize features that:

1. High Impact on Retention

- Features that directly reduce churn risk
- Example: Real-time empathy alerts

2. High Confidence through Validation

- Validated through pilot data

- Example: Post-release spike detection (proven need)

3. Reasonable Effort

- Leverage existing infrastructure
- Example: Slack alerts (integration already exists)

Specific Example:

Feature A: Agent-level performance trends

- Impact: 8 (helps CX managers coach)
- Confidence: 70% (pilots show managers want this)
- Effort: 5 (requires new aggregation logic)
- Score: $(8 \times 0.7) / 5 = 1.12$

Feature B: Voice call transcription integration

- Impact: 9 (huge volume of voice interactions)
- Confidence: 60% (transcription quality unknown)
- Effort: 8 (complex integration)
- Score: $(9 \times 0.6) / 8 = 0.68$

Decision: Prioritize Feature A

Also consider:

- Strategic alignment (what does leadership want?)
- Dependencies (what blocks other features?)
- Technical debt (what makes future work easier?)
- "

Q: "You're 3 months in. CSAT hasn't improved. What do you do?"

Your Answer:

"First, I wouldn't panic. CSAT is a lagging indicator. Here's my approach:

Step 1: Check Leading Indicators (Day 1-2)

- Has Conversation Quality Score improved?
- Is FCR up? ART down?
- Are agents using the insights?

If leading indicators ARE improving:

- CSAT lag is normal (30-60 day delay)
- Continue course, communicate lag to stakeholders

- Set expectations: 'We'll see CSAT improvement in Month 4-5'

If leading indicators AREN'T improving:

Step 2: Diagnose the Problem (Week 1)

Hypothesis A: Adoption Issue

- Check: Dashboard usage stats
- If low: User interviews to understand barriers
- Fix: Training, UX improvements, change management

Hypothesis B: Data Quality Issue

- Check: AI accuracy vs human evaluators
- If poor: Re-validate model, refine prompts
- Fix: Retrain on recent data

Hypothesis C: Wrong Dimensions

- Check: Correlation analysis (do our dimensions predict CSAT?)
- If weak: User research to identify missing factors
- Fix: Add/modify dimensions

Hypothesis D: External Factors

- Check: Product quality, pricing changes, competition
- If external: Adjust expectations, focus on relative improvement
- Fix: Work with Product team on systemic issues

Step 3: Communicate Transparently (Ongoing)

- Share findings with stakeholders weekly
- Show what you're learning
- Adjust strategy based on data

Step 4: Pivot if Necessary (Week 2-3)

- If fundamental assumption wrong: Admit it, propose new approach
- Better to pivot fast than stick to failing plan

Key Principle: 'Trust the process, but validate assumptions.' Leading indicators should improve first. If not, investigate systematically."

Q: "Product team says your insights are too detailed. What now?"

Your Answer:

"This is a great signal - it means they're engaging with the product! Here's how I'd respond:

Step 1: Understand the Real Problem (Day 1)

- Schedule 30-min conversation
- Ask: 'What decisions are you trying to make?'
- Understand: 'What's overwhelming you?'
- Clarify: 'What would make this more useful?'

Step 2: Create Persona-Specific Views (Week 1)

For Executives:

- Top 3 insights only
- Business impact highlighted
- Weekly email digest

For Product Managers:

- Filtered by product area
- Prioritized by impact × volume
- Monthly report with recommended actions

For Engineers:

- Deep-dive access available
- But default view is high-level

Step 3: Add Filtering & Customization (Week 2-3)

- User-configurable dashboard
- Save favorite views
- Alert preferences (how much do you want to hear from us?)

Step 4: Implement Progressive Disclosure (Week 3-4)

- Show summary first
- 'Dig deeper' option for details
- Breadcrumb navigation

Example Redesign:

Before: Dashboard with 6 heatmaps, 10 charts, 50 data points

After:

- Hero metric: 'Top Issue This Week'

- 3 tiles: Urgent, Important, FYI
- Click to expand for full analysis

Step 5: Measure Success

- Increased usage frequency
- Positive feedback in follow-up
- Insights being acted upon (track actions taken)

Key Insight: 'More data doesn't mean more insight.' Job is to reduce cognitive load, not increase it."

SentiSum-Specific Questions

Q: "Why SentiSum? Why this role?"

Your Answer:

"Three reasons:

1. Product Mission Resonates

I'm passionate about products that solve hard, important problems. Customer churn is existential for businesses, but most VoC tools are just survey dashboards. SentiSum is different - it's AI-native, real-time, and actually drives action. That's rare and exciting.

2. Role is Perfect Blend

This role combines my product management experience with my growing interest in AI. I've built MVPs and validated them with users (wellness app at Dentira). Now I want to apply that user-centered approach to AI products specifically. This role is that perfect overlap.

3. Team and Culture Fit

From the JD: 'Products should adapt to users, not vice versa' - I couldn't agree more. Your emphasis on shipping fast, testing with real users, and iterating resonates with how I work. Small team, high impact, founder-level exposure - that's where I thrive.

Specific Interest in Kyo:

The AI assistant concept is fascinating. Moving from passive dashboards to proactive intelligence is the future of enterprise software. I'd love to contribute to that evolution - making AI feel less like a tool and more like a thoughtful team member.

What I Bring:

I'm not an AI expert (yet!), but I'm a fast learner with strong product fundamentals. I obsess over user problems, I validate with data, and I execute quickly. For this stage of SentiSum, I think that product builder mentality is exactly what you need."

Q: "What do you know about our product?"

Your Answer:

"I've done my homework! Here's what I understand:

Core Product:

- AI-native VoC analytics platform
- Unified feedback from 100+ sources (tickets, reviews, calls, surveys, social)
- Real-time analysis and anomaly detection

Key Differentiator: Kyo

- Acts as 'Customer Second Brain'
- Conversational interface - ask questions, get answers
- Proactive insights (doesn't wait for you to look)
- Role-specific intelligence

Technical Capabilities:

- Speech analytics for voice calls
- Multi-language support
- Custom AI models per customer
- Root cause analysis (not just sentiment)

Target Customers:

- Mid-market to enterprise
- High customer interaction volume
- Support, CX, Product, Retention teams
- Companies serious about reducing churn

Recent Focus:

- Anomaly detection improvements (catch issues in hours, not days)
- Cross-channel correlation (trace issue from review → ticket → social)
- Real-time alerting

What I'm Curious About:

- How do customers typically first engage with Kyo? Queries or proactive insights?
- What's the split between support-heavy vs survey-heavy customers?
- What's the biggest technical challenge right now?"

7. Final Preparation Checklist

Day Before Interview

☐ Re-read All Your Assignment Documents

- FinServeCo-AI-Quality-Framework.pdf
- JSON-Output-Schema.pdf
- FinServeCo-Project-Plan.pdf
- Customer-Support-Ticket-Analysis-Prompt.pdf

☐ Review SentiSum Online Presence

- Website (senti-sum.com)
- LinkedIn company page (recent posts)
- Any recent press or blog articles

☐ Prepare Your Stories

- 3 examples from past work (STAR format)
- 1 failure/learning story
- 1 collaboration challenge story

☐ Write Down Your Questions

- 5 thoughtful questions for them
- Mix of role, product, and culture questions

☐ Practice Key Explanations

- Explain CoT prompting in 60 seconds
- Draw your KPI tree from memory
- Summarize your assignment in 3 minutes

30 Minutes Before Interview

☐ Technical Setup

- Test camera, mic, internet
- Close all other applications
- Have assignment PDFs open on second screen
- Have blank document ready for notes

☐ Physical Setup

- Good lighting on your face
- Neutral background

- Water nearby
- Phone on silent

☐ **Mental Setup**

- Review your 3 key strengths
- Deep breathing (5 breaths)
- Remind yourself: "I've prepared well. I've got this."

During Interview

Opening (First 5 min)

- ☐ Smile and make eye contact (camera!)
- ☐ Thank them for the opportunity
- ☐ Express genuine excitement about the role

Assignment Walkthrough (30-40 min)

- ☐ Start with the problem, not your solution
- ☐ Use the 'Problem → Framework → Validation' structure
- ☐ Draw diagrams if helpful
- ☐ Pause to check if they want more detail

Deep Dive Questions (25-30 min)

- ☐ Listen carefully to the full question
- ☐ Ask for clarification if needed
- ☐ Structure your answer (1, 2, 3...)
- ☐ Give specific examples
- ☐ Connect back to business impact

Your Questions (5-10 min)

- ☐ Ask your prepared questions
- ☐ Listen actively to their answers
- ☐ Follow up with related questions
- ☐ Show curiosity about the details

Closing (Last 2 min)

- ☐ Thank them again
- ☐ Reiterate your interest
- ☐ Ask about next steps
- ☐ End on a positive note

After Interview

Same Day:

- ☐ Send thank-you email within 6 hours
- ☐ Reference specific conversation points
- ☐ Reiterate 1-2 key strengths you bring
- ☐ Express continued enthusiasm

Within 24 Hours:

- ☐ Reflect on what went well
- ☐ Note any questions you struggled with
- ☐ Prepare for potential follow-up rounds

Key Reminders

You've Got This Because:

- ✓ Your assignment is excellent and well-thought-out
- ✓ You understand both AI and product fundamentals
- ✓ You have relevant experience (MVP development, user research)
- ✓ You're prepared (you've read this entire guide!)
- ✓ You're genuinely interested in the problem space

If You Get Stuck:

- Take a breath
- Ask for clarification
- It's okay to say "That's a great question, let me think..."
- Be honest if you don't know - show how you'd figure it out

Remember:

They're not looking for perfection. They're looking for:

- **Clear thinking** - can you structure problems?
- **User empathy** - do you care about solving real problems?
- **Execution ability** - can you ship things?
- **Learning mindset** - can you grow into the role?

You have all of these. Trust yourself.

Additional Resources Created

CSV Files for Quick Reference:

1. AI_Concepts_QA.csv - 8 core concepts
2. AI_Terminology_Guide.csv - 20 key terms
3. PM_Concepts_Reference.csv - 15 frameworks

4. Assignment_Questions_Answers.csv - 13 likely questions

5. FinServeCo_KPI_Tree.csv - Complete hierarchy

Master Guide:

- Interview_Prep_Master_Guide.txt - Complete checklist

Good luck, Bhavya! 🍀

You're well-prepared, your work is strong, and you're going to do great.

Be confident. Be curious. Be yourself.